SPE-229240-MS

Agentic AI Revolution Across O&G Value Streams: New Strategy Through ENERGYai Examples

Hans Jessen, ADNOC, UAE; Mike Roshchin, Inception/G42, UAE

This paper was prepared for presentation at the ADIPEC held in Abu Dhabi, UAE, 3 – 6 November 2025.

Abstract

Artificial Intelligence (AI) is rapidly transforming the oil and gas (O&G) industry, with the emergence of Agentic AI marking a major leap forward. Unlike traditional rule-based analytics or static machine learning models, Agentic AI refers to autonomous AI agents capable of perceiving their environment, learning continuously, and taking independent actions to achieve defined objectives — all with minimal human intervention. These agents are not limited to isolated tasks; instead, they can operate across the entire O&G value chain, from upstream exploration to downstream refining and distribution, making decisions and executing operations in real time.

A pioneering example of this shift is ADNOC's recently launched ENERGYai platform — the industry's first large-scale deployment of Agentic AI. ENERGYai integrates specialized autonomous agents across ADNOC's operations, performing tasks such as seismic interpretation, reservoir management, production optimization, and refining process monitoring. These agents combine domain-specific intelligence with the reasoning capabilities of large language models (LLMs), such as OpenAI's GPT, to enhance decision-making and workflow integration.

By leveraging commoditized foundation models and AI orchestration frameworks, ENERGYai demonstrates how O&G companies can achieve significant operational efficiencies, reduce downtime, and improve safety. Agile AI-driven teams are increasingly able to outperform traditional legacy systems by responding to real-time data and conditions, creating a competitive edge through automation and intelligent system coordination.

This white paper explores the transformational potential of Agentic AI across upstream, midstream, and downstream segments. Using ENERGYai as a real-world case study, we delve into its applications in seismic data analysis, production forecasting, and autonomous process control. We also examine how combining LLMs with autonomous agents accelerates AI adoption and enables broader organizational agility.

However, the adoption of Agentic AI is not without challenges. Data quality, QA/QC protocols, inconsistent data standards (e.g., the need for alignment with the OSDU Data Platform), and concerns around trust and explainability remain key hurdles. This paper addresses how leading O&G players are navigating these issues to unlock the full value of autonomous agents while ensuring operational transparency, governance, and stakeholder trust.

In summary, Agentic AI represents a transformative shift in the digitalization of O&G, offering a scalable, intelligent, and adaptive future for energy operations.

AI Agents vs Legacy Systems

One striking aspect of agentic AI is how it empowers **agile teams** to outperform traditional operations and legacy systems. In the past, optimizing an oilfield or running a refinery relied heavily on large teams of specialists, siloed legacy software, and time-consuming manual workflows. Decisions (like field development planning or production scheduling) often took months or years of analysis using legacy tools. Today, a small, cross-functional team armed with AI agents can iterate on solutions in days – a dramatic boost in agility. For example, ADNOC noted that with ENERGYai, building detailed geological models for carbon storage projects could be accelerated by up to **75%**, compressing workflows that took months into a matter of days[1][3]. Similarly, development plans that traditionally required 1–2 years of studies can potentially be generated in a few weeks by an AI-powered system that explores countless scenarios in parallel. This ability to rapidly simulate and evaluate multiple options allows teams to respond faster to new information or changing objectives, something legacy processes struggled with.

Agile teams leveraging agentic AI function like **"digital start-ups" within the enterprise**, rapidly prototyping and deploying new models or strategies. Because foundational AI capabilities (like LLMs and cloud computing) are now widely available, teams do not need to build everything from scratch – they can assemble solutions quickly by combining commoditized components with domain expertise. This enables even smaller operators or new entrants to leapfrog legacy systems. In production operations, for instance, an agile team might deploy a suite of AI agents (one for well optimization, one for pump maintenance, another for logistics) that together outperform a decades-old monolithic SCADA system which is rigid and cannot learn. The AI agents continuously adapt and learn from data, whereas legacy systems often rely on static rules set by engineers long ago. In practice, we see **faster decision loops**: an AI agent detects a problem or opportunity, recommends an adjustment, and the human team can implement or oversee it immediately – rather than waiting for a monthly engineering meeting or a manual analysis report. Over time, as trust in the AI grows, many of these decisions can be automated entirely.

Another advantage is how **agentic AI breaks down silos** that plague legacy operations. Traditional O&G organizations often had separate teams (geoscience, drilling, reservoir engineering, operations) each using their own software and datasets. Coordinating changes (e.g. adjusting a production plan) across these silos was slow. An AI agent, however, can ingest data and objectives spanning all these domains and produce an integrated action plan. This not only speeds up execution but often yields better outcomes by considering the full system rather than sub-optimizing parts. Agile teams, freed from mundane data-crunching (since AI handles that), can focus on higher-level strategy and creative problem-solving. In essence, **agentic AI serves as a force-multiplier for human teams** – handling routine analyses and first-pass decisions at machine speed, so the team can concentrate on innovation and oversight[4]. Early adopters report significant efficiency gains: for example, at Saudi Aramco's digitized Khurais field, AI and automation measures delivered a 15% increase in oil production and doubled the speed of troubleshooting issues[4]. Such outcomes underscore that a small team armed with the right AI tools can achieve in months what a legacy operation might take years to do, all while maintaining safety and reliability.

Foundation Models as Accelerators for AI Adoption

A key enabler of this AI revolution in O&G is the rise of **commoditized foundational models**, especially large language models. These are massive pre-trained AI models (like GPT-4, etc.) that can be adapted to a wide range of tasks – from understanding unstructured text to reasoning over complex problems. In the past, deploying advanced AI in a specialized industry meant developing custom models for each task at great cost and effort. Today, however, baseline models are available via cloud APIs or open-source communities, dramatically lowering the barrier to entry. Industry leaders describe foundation models as increasingly akin to a utility or commodity – widely accessible and cost-effective to use[5]. This commoditization means the competitive advantage shifts from *owning* a model to *how you use it*. For O&G companies, this is liberating: they can focus on applying AI to domain problems (seismic, drilling, production, etc.) without needing an army of PhD researchers to invent AI from scratch.

In fact, **agentic AI relies heavily on foundation models as the "brain"** powering the agent's cognitive abilities[6]. An autonomous agent typically needs to perceive inputs (which could be images, sensor data, or text), reason about what it means, and decide on actions. Foundation models provide advanced capabilities like natural language understanding (to read technical reports or instructions), pattern recognition, and even planning via their trained knowledge. By plugging these general models into energy-specific data, companies can get powerful results with little training data of their own. The ENERGYai initiative explicitly integrates OpenAI's models along with industry data platforms (like OSDU) to leverage best-in-class AI without reinventing the wheel[1][3].

The result is a significant **acceleration in AI solution development**. As one analysis noted, the commoditization of foundation models is expected to *"accelerate the development of domain-specific AI agents"*, giving enterprises more choices and better quality AI solutions, faster[6]. Essentially, with cheaper and more accessible base models, an upstream team can develop an AI agent to, say, optimize gas lift in wells in a matter of weeks by fine-tuning an existing model – whereas a few years ago this might have been a multi-year custom AI project. The **democratization of AI** is underway: cost reductions, open-source models, and greater explainability are leading to a surge in adoption of agentic AI across industries[6]. Even mid-sized operators or service companies can now harness AI, not just the super-majors with big R&D budgets. Analysts predict this trend will result in agentic AI becoming far more widespread, as *the barriers to entry have been brought down significantly*, allowing companies to embark on AI initiatives "without a heavy financial luggage"[6]. In practical terms, we see many vendors offering pre-trained energy-specific AI modules (for example, fault detection algorithms, drilling optimization agents) built atop general models. Companies can mix and match these to rapidly assemble solutions tailored to their assets. It's worth noting that **explainability and transparency of models are improving as well** – newer foundation models (and techniques like prompt engineering or model fusion) are making it easier to interpret AI decisions[6][6]. This is critical for trust (discussed next). In summary, commoditized LLMs and foundation models act as a catalyst, **supercharging the adoption of agentic AI** in oil and gas by making advanced capabilities readily available and integration friendly.

Key Challenges: Data Quality, Standardization, and Trust

Despite the impressive progress and potential of agentic AI in O&G, several challenges must be addressed to fully realize its benefits. **Data Quality and QA/QC:** Oil and gas operations generate colossal amounts of data – from seismic traces and well logs to sensor readings on every piece of equipment. However, *raw data quality is often a limiting factor*. No AI can produce reliable insights from bad data ("garbage in, garbage out"). Many organizations find that their historical datasets contain errors, inconsistencies, or biases – for example, sensor calibrations drift over time, different fields use different naming conventions, or older records might be incomplete. This makes rigorous **data QA/QC (Quality Assurance/Quality Control)** an essential foundation for any AI project. Before deploying agents, leading companies invest in cleaning and curating their data: outlier detection, data reconciliation, and establishing single sources of truth for key metrics. In sensitive domains like well control or pipeline integrity, *validation of data and AI outputs is non-negotiable* – engineers will often have the AI's recommendations cross-checked against physics-based models or human expertise as a quality gate. As a result, many autonomous agents are initially introduced in a **decision support role** (providing recommendations) rather than fully automatic control, until sufficient confidence and validation have been achieved. This phased approach is part of QA/QC to ensure that when the AI does act autonomously, it is doing the right thing.

Lack of Standardized Data & the OSDU Initiative

A related challenge is the historical lack of data standardization in O&G. Different companies – even different departments – have stored data in proprietary formats and siloed databases. Subsurface data might reside in one system, production data in another, and they often are not easily compatible. This fragmentation hinders the development of AI that needs to draw insights into all available information. Industry leaders recognize this and have launched the **Open Subsurface Data Universe (OSDU)** initiative, a collaborative effort to create a common, open data platform and format for energy data. The OSDU platform is **cloud-native, technology-agnostic, and designed for the upstream industry**, aiming to liberate data from silos and encourage seamless integration of applications[7]. By adopting OSDU standards, data from seismic surveys, wells, production, etc., can be stored in a single framework with common semantics. This facilitates AI/ML applications. As one Wipro whitepaper noted, *valuable subsurface data is often "trapped in incompatible proprietary formats," hindering the industry's evolution, and the absence of centralized, standardized, high-quality datasets pose a significant challenge*. Embracing OSDU (and similar standards) is seen as the solution to this obstacle. With a unified data platform, an AI agent can quickly access all relevant data types, and developers can build AI models that work across datasets without custom connectors for each source. The OSDU also promotes an ecosystem of interoperable tools, meaning an agent developed by one vendor can potentially plug into another company's OSDU-compliant data store – reducing vendor lock-in and fostering innovation. However, moving towards standardization is a non-trivial effort. Companies must migrate legacy data, ensure metadata quality, and often re-think workflows to align with the new platform. This is an ongoing journey (the OSDU Forum continues to refine standards), but the momentum is strong because the payoff – robust AI-driven insights at scale – is so compelling[7].

Trust and Transparency in AI Agents

The most significant challenge is not technical but human: *building trust in autonomous AI agents*. In an industry where safety and reliability are paramount, there is understandable caution about handing over critical decisions to an algorithm. Operators and engineers need to trust that an AI agent's recommendations or actions are correct, safe, and consider all constraints. A lack of trust can lead to user pushback or underutilization of AI capabilities – indeed, *concerns around the lack of trust in AI outputs have been a major factor slowing down fully autonomous operations*[6]. To address this, AI systems in O&G are being designed with **transparency and explainability** in mind. Rather than a "black box" giving cryptic suggestions, modern AI platforms attempt to provide reason codes or highlight the data driving a recommendation. For example, an AI agent that suggests increasing the injection rate in a well might also show which sensor trends or simulation results support that suggestion. This helps experts verify that the AI's logic aligns with engineering principles. Additionally, companies often maintain a "human-in-the-loop" for high-impact decisions: the AI can act autonomously for routine optimizations, but if an action crosses certain risk thresholds, it flags a human supervisor to approve. Over time, as the AI proves itself by consistently making good decisions (even better than human operators in some cases), trust grows, and more autonomy can be granted.

Another approach to foster trust is **extensive field testing and phased rollout**. For instance, during the development of ENERGYai, ADNOC is conducting a three-year program testing the agents on real-world datasets in specific domains (seismic analysis, CO_2 storage modeling, etc.)[1]. By demonstrating concrete improvements (like the 75% faster geological modeling, or significant accuracy gains) and refining the system with user feedback, they build credibility among the workforces. It is also important to involve end-users (geoscientists, engineers, plant operators) in the development loop – through user acceptance testing and iterative improvement – so they feel ownership of the AI tools rather than seeing them as an imposed technology. **Transparency in model behavior** is receiving increased attention as well. New foundation models are being developed with *explainable AI* techniques; for example, some provide explanations for their outputs or have built-in traceability for decisions[6][6]. An example from the AI research community is the emergence of models that can highlight which data points most influenced their decision, or that can answer questions about why they made a recommendation. Such features, though still maturing, can help in a domain like O&G where experts will challenge and probe any suggestion.

Finally, **governance and ethical guidelines** play a role. Companies are establishing AI governance frameworks that define the boundaries for autonomous agents – e.g. an agent cannot override safety interlocks or must respect certain operating limits at all times – ensuring that no matter how intelligent the AI, it adheres to the company's safety and compliance standards. Regulators too are beginning to pay attention to industrial AI, and while specific regulations for AI in O&G are still nascent, we can expect requirements for auditability of AI decisions in critical operations. All these measures contribute to making AI agent outputs more trustworthy and transparent. When operators see that an autonomous system consistently makes the right call *and* they understand why confidence builds. Over time, the culture shifts from one of skepticism to one of seeing AI as a reliable partner. Indeed, agentic AI is often positioned not as a replacement for human expertise but as a **collaborative teammate** – performing heavy computational lifting and routine tasks, while humans provide oversight, domain judgment, and handle novel situations. With this approach, trust can be earned gradually, unlocking the full potential of autonomous agents safely.

Upstream: Intelligent Exploration and Production

Upstream exploration and production stand to gain enormously from agentic AI. **Exploration & Seismic Interpretation:** AI agents can analyze vast geological datasets (seismic surveys, well logs, etc.) far faster and more accurately than teams of geoscientists. In fact, ADNOC's trial of the ENERGYai agents demonstrated a *70% improvement in accuracy* on key aspects of seismic interpretation, significantly accelerating the discovery of new hydrocarbon prospects while reducing costs[4]. Tasks like seismic fault detection or reservoir structure mapping that once took months of manual analysis can now be completed in days using AI – as Sultan Al Jaber (ADNOC's CEO) highlighted, agentic AI *"will speed up seismic surveys from months to days"*[2]. These gains come from AI's ability to sift petabytes of subsurface data and recognize subtle patterns that humans might miss. Additionally, AI-driven reservoir modeling allows engineers to continuously update and refine subsurface models. Such models provide high-fidelity visualizations of the subsurface, helping pinpoint optimal drilling targets and better field development plans.

Autonomous Drilling & Production Optimization

Once resources are identified, AI agents improve drilling and well operations. Nearly 29% of oil companies globally are already piloting agentic AI in drilling control[4]. These autonomous systems act as copilots on the rig, continuously adjusting drilling parameters (weight on bit, mud pump rates, etc.) to optimize performance and safety. For example, an AI agent can automatically modify drilling mud weight in real-time to prevent a well blowout, learning from sensor feedback (pressure, vibration) and react instantaneously to avoid incidents. By responding to early signs of instability faster than any human could, such agents minimize costly downtime and accidents. Beyond drilling, AI agents manage day-to-day production by tuning well settings (choke valves, gas injection rates, artificial lift settings) to maximize output while preventing problems like water breakthrough or sand influx. ADNOC's intelligent reservoir management system **AR360** exemplifies this autonomous optimization: AR360 integrates data from dozens of wells and uses AI to **autonomously balance new drilling with optimization of existing wells**, continually reallocating resources to maximize total field production. In effect, AI agents serve as tireless production engineers, tweaking operations 24/7 to squeeze more output from the reservoir without breaching safety or reservoir integrity constraints.

Predictive Maintenance of Upstream Assets

Agentic AI also excels at equipment monitoring and maintenance – a critical need in upstream, where unplanned downtime can cost millions. By analyzing sensor data from rigs, pumps, and compressors, AI agents can detect *early warning signs* of equipment failure that humans might overlook. For instance, vibration patterns that hint at a drill-bit bearing wearing out, or pressure fluctuations that signal a pump leak, can be caught by AI well before a catastrophic failure occurs[4]. This predictive maintenance approach extends asset life and avoids outages. Industry studies show that AI-driven predictive maintenance programs have improved production uptime by about **27%** on average. ADNOC recently deployed a centralized predictive analytics system (CPAD) which continuously monitors upstream operations, diagnoses anomalies from historical and real-time data, and proactively recommends maintenance actions – thereby significantly reducing unplanned downtime[4]. Not only does this improve safety (e.g. catching a small casing corrosion before it causes a blowout or spill), but it also yields environmental benefits by preventing incidents. In summary, **upstream agentic AI delivers a step-change in performance** – faster and more precise exploration, safer and more efficient drilling, and self-optimizing production operations. By reducing manual effort and human error, these systems boost output while cutting costs and operational risks. This is crucial for operators as fields mature, and easy oil is depleted; AI helps maintain high productivity and enables more sustainable development of new resources.

Midstream: Smarter Transport and Logistics

Midstream operations – pipelines, storage terminals, and transportation networks – benefit from autonomous AI agents that monitor infrastructure in real time and optimize the flow of hydrocarbons. **Pipeline Monitoring & Leak Detection:** O&G companies oversee thousands of kilometers of pipelines, many aging. AI agents are being deployed as round-the-clock pipeline sentinels, analyzing pressure sensors, flow meters, acoustic data, and even satellite imagery to catch anomalies far earlier than traditional methods. Subtle pressure drops or vibration patterns that might indicate a small leak, corrosion hotspot, or ground movement can be picked up by machine learning models long before a human operator might notice. In fact, roughly *23% of O&G operators are already testing AI agents for pipeline monitoring*, using sensor and satellite data to proactively detect issues before they escalate[4].

Early leak detection not only averts environmental damage but also reduces product loss and downtime, delivering both safety and economic benefits. AI can further predict integrity issues – forecasting corrosion growth or fatigue cracks – so that pipeline segments can be repaired or replaced just in time, rather than after a failure. This data-driven predictive approach maximizes the operational life of pipeline assets.

Autonomous Flow Control and Logistics

Beyond surveillance, Agentic AI can actively manage midstream networks to optimize throughput. These AI "dispatchers" ingest data on production rates, pipeline pressures, tank storage levels, shipping schedules and even market demand, and then make real-time decisions on how to route oil and gas most efficiently. For instance, an AI agent might dynamically reroute crude oil flows across a network of pipelines and storage hubs to avoid bottlenecks or reduce energy costs (pumping power), all while respecting safety limits on pressures[4]. This kind of **autonomous midstream optimization** can reduce bottlenecks and ensure customers get the right products at the right time with minimal waste.

Midstream Asset Maintenance

Like upstream, midstream infrastructure benefits from predictive maintenance by AI. Compressor stations, pumps, LNG terminals, and large storage tanks operate continuously, so a single unplanned outage can be extremely costly (industry analysts note that a **single day of unplanned downtime** of a major pipeline or refinery can incur losses over $1 million). AI models learn normal operating signatures (vibration, temperature, pressure) of rotating equipment and can alert operators to deviations that suggest incipient bearing failures, seal leaks, or valve issues. By scheduling maintenance at optimal times (during low throughput or planned downtime), companies avoid catastrophic breakdowns and extend equipment life. ADNOC's CPAD system, mentioned earlier, not only covers upstream but also diagnoses anomalies in midstream and downstream facilities, recommending maintenance to ensure smooth operations and minimal disruption. Additionally, AI-powered robots and drones are joining the midstream workforce – for example, Many O&G companies use already today camera-equipped drones and ground robots to routinely inspect remote pipelines for faults, guided by AI vision algorithms. These autonomous or semi-autonomous inspection agents make maintenance safer (reducing the need for personnel in hazardous or hard-to-reach areas) and more efficient.

Logistics and Supply Chain Optimization

Midstream value also comes from smarter logistics. AI agents can optimize vessel scheduling, trucking routes, and storage management. For instance, ADNOC implemented an **Integrated Logistics Management System (ILMS)** powered by AI to coordinate the movement of petroleum products across its network[4]. This system automatically schedules and routes deliveries (e.g. dispatching oil tankers or gas trucks) in an optimal way, considering factors like inventory levels at depots, shipping constraints, and customer demand. By providing planners with data-driven options for scheduling, such AI solutions reduce waiting times and ensure high asset utilization in midstream logistics. In summary, **midstream agentic AI improves reliability, efficiency, and safety** of the transport and storage network. Real-time monitoring prevents leaks and failures, autonomous control maximizes flow efficiency, and predictive maintenance plus optimized logistics minimize downtime and costs. These enhancements not only lower operating costs but also strengthen environmental stewardship by preventing incidents – a crucial benefit as operators modernize decades-old infrastructure for the future.

Downstream: Refining and Distribution Automation

Downstream operations – refining crude into products and delivering them to end-users – are complex, margin-sensitive, and asset-intensive. Agentic AI is enabling **"smart refineries"** and distribution networks that run with far greater precision, yield, and reliability than legacy setups. **Refinery Process Optimization:** Modern refineries have thousands of control points and must continuously adjust to changing crude qualities, product specs, and energy costs. AI agents can act as expert process operators, automatically tweaking settings to maximize output and energy efficiency in real time. Already, many refiners use advanced analytics to optimize yields (for example, adjusting temperatures in distillation columns or catalytic crackers). Agentic AI takes this further by *learning from decades of historical process data and real-time sensor streams to find optimal operating points* that even veteran engineers might overlook.

Product Quality Control and Yield Improvement

AI agents excel at recognizing patterns in process data that relate to product quality. In refining, they can forecast properties like gasoline octane or diesel sulfur content based on current conditions, and then proactively adjust process parameters to meet specifications[4]. This reduces off-spec production and the need for relending or reprocessing, directly saving costs. Moreover, by analyzing subtle correlations, AI may discover ways to slightly boost the yield of high-value products from each barrel of oil – effectively squeezing more value out of the same input. For example, an "AI process engineer" agent might identify that a small tweak in a unit's temperature and pressure setpoints yields a higher fraction of petrochemical feedstock without harming fuel quality. Early deployments of AI in downstream report about a **26% improvement in asset utilization on average**, reflecting more efficient use of refining capacity and better matching of output to demand.

Predictive Maintenance in Refineries

Refineries and petrochemical plants contain tens of thousands of pieces of equipment (pumps, heat exchangers, valves, furnaces). A failure of any critical component can force a partial or full shutdown, which is extremely costly. Autonomous maintenance agents significantly reduce this risk by monitoring equipment health and performance trends. They ingest sensor readings – temperatures, vibrations, ultrasonic thickness gauges, etc. – and detect patterns that precede failures, issuing early warnings or even automatically generating work orders for parts replacement. This is vital for older refineries in the region that have been operating for decades. At the Abqaiq plant mentioned above, advanced analytics and AI predictive models now help engineers **anticipate potential system failures well in advance**, allowing maintenance to be scheduled during planned downtime. By catching issues early (e.g. a slight vibration increase indicating misalignment, or a slow decline in heat exchanger performance indicating fouling), the plant minimizes unplanned outages despite its age[4]. Predictive maintenance can accurately pinpoint early signs of deterioration like corrosion or wear, so maintenance can be done proactively without disrupting operations. The economic payoff is clear: avoiding just one major unplanned refinery outage (often >$1M per day in losses) easily justifies enterprise-wide AI maintenance programs. Companies are extending these AI maintenance strategies not just to refineries but across their downstream fleets of equipment to keep utilization high.

Conclusion

Agentic AI is ushering in a transformative era for the O&G industry. Across upstream exploration, midstream logistics, and downstream refining, we see autonomous AI agents delivering improvements in speed, accuracy, efficiency, and safety that were scarcely imaginable a decade ago. Real-world deployments like ADNOC's ENERGYai highlight how integrating specialized AI agents with decades of domain data can yield step-change outcomes – from seismic analysis done in days instead of months, to production forecasts that are far more accurate[2]. Agile teams empowered by these AI tools are outperforming legacy systems, as continuous machine-driven optimization replaces slow manual processes. The **commoditization of foundational AI models** has accelerated this trend, making advanced capabilities accessible to all and allowing the industry to focus on creative applications rather than reinventing AI technology from scratch[6].

Yet, as with any technological revolution, success depends on addressing the challenges. High-quality data and standardized platforms (like OSDU) are the fuel for AI – ensuring that agents have reliable information and a complete picture of operations. Trust and transparency are the license to operate; the workforce and stakeholders must have confidence in autonomous decisions, built through explainable systems, rigorous QA/QC, and human oversight where needed. The O&G industry's journey with agentic AI is still in its early chapters. Over the next few years, as more pilot projects mature into scaled solutions, we can expect the **global energy sector to set new benchmarks** for AI-enabled performance[3]. Upstream discovery might become faster and more efficient, enabling more sustainable resource development. Midstream networks will run closer to peak efficiency with AI orchestrating flows and maintenance. Downstream plants could achieve optimum yields with minimal emissions through self-tuning processes. Importantly, these improvements also align with broader goals like safety excellence and emissions reduction, demonstrating that AI can drive both profit and sustainability in tandem[4].

In closing, agentic AI is not a magic button but a powerful toolkit – one that, when applied with domain knowledge, data discipline, and ethical guardrails, can unlock enormous value in oil and gas operations. Companies that embrace this technology and cultivate trust in it stand to gain a competitive edge in an industry challenged by volatile markets and energy transition. Those who hesitate may find themselves outpaced by more agile, AI-enabled players. The evidence so far suggests that the future O&G enterprise will be one where **autonomous AI agents work together with human experts**, delivering a level of operational agility and insight that sets new standards for what is possible in the energy industry. The transformation has begun, and it is both exciting and imperative for industry leaders to navigate it with clarity, embracing innovation while upholding the principles of safety, transparency, and collaboration that underpin the sector's license to operate.

References

1. ADNOC and AIQ Developing First-of-a-Kind Agentic AI Solution for Global Energy Transformation – ADNOC https://www.adnoc.ae/en/news-and-media/press-releases/2024/adnoc-and-aiq-developing-first-of-a-kind-agentic-ai-solution-for-global-energy-transformation
2. UAE'S ADNOC to deploy autonomous AI in the energy sector for the first time | Reuters https://www.reuters.com/business/energy/uaes-adnoc-deploy-autonomous-ai-energy-sector-first-time-2024-11-04/
3. ADNOC and AIQ launch agentic AI solution for global energy transformation https://worldoil.com/news/2024/11/4/adnoc-and-aiq-launch-agentic-ai-solution-for-global-energy-transformation/
4. The Future Role of Autonomous AI Agents in the Energy Value Chain (UAE & Saudi Arabia) https://www.linkedin.com/pulse/future-role-autonomous-ai-agents-energy-value-chain-uae-kumar-fgwaf
5. Foundational models as commodity? https://www.v4.ventures/foundational-models-as-commodity
6. The DeepSeek Effect: Impact of Foundation Model Commoditization on Agentic AI Adoption https://www.ema.co/blog/agentic-ai/the-deepseek-effect-impact-of-foundation-model-commoditization-on-agentic-ai-adoption
7. How AI Will Supercharge the OSDU – Wipro https://www.wipro.com/oil-and-gas/how-ai-will-supercharge-the-osdu/