



Society of Petroleum Engineers

SPE-229446-MS

Agentic AI: A Tool-Guided Approach to Achieving Precision in Engineering Numerical Operations

Mallikarjuna Rao Gudipati, SLB, Data Hub, Pune, India

Copyright 2025, Society of Petroleum Engineers DOI [10.2118/229446-MS](https://doi.org/10.2118/229446-MS)

This paper was prepared for presentation at the ADIPEC held in Abu Dhabi, UAE, 3 – 6 November 2025.

This paper was selected for presentation by an SPE program committee following review of information contained in an abstract submitted by the author(s). Contents of the paper have not been reviewed by the Society of Petroleum Engineers and are subject to correction by the author(s). The material does not necessarily reflect any position of the Society of Petroleum Engineers, its officers, or members. Electronic reproduction, distribution, or storage of any part of this paper without the written consent of the Society of Petroleum Engineers is prohibited. Permission to reproduce in print is restricted to an abstract of not more than 300 words; illustrations may not be copied. The abstract must contain conspicuous acknowledgment of SPE copyright.

Abstract

This paper introduces an agentic AI architecture that augments Large Language Models (LLMs) with tool-guided workflows to achieve consistent and dependable mathematical extraction and computation. By guiding LLMs through structured, agentic workflows, the proposed system addresses their inherent underperformance in mathematical operations and numerical consistency, despite strong contextual understanding. The system decomposes tasks into a sequence of modular, autonomous agents—each assigned a distinct function—ensuring high-precision outcomes in data-driven engineering applications.

The architecture operates through a well-defined process: a Document Agent ingests and preprocesses unstructured inputs (e.g., PDF engineering reports); an Extraction Agent utilizes an LLM to identify and extract structured information (e.g., acid volumes, fluid stages, depth ranges); a Validation Agent performs schema-based validation to ensure data integrity; a Computation Agent executes deterministic calculations on validated data (e.g., summing acid volumes, determining total Viscoelastic Diverting Acid volumes, and calculating depth intervals); and a Reporting Agent formats and outputs the results as structured tables or JSON.

The framework was tested using acid job reports, successfully parsing, validating, and processing examples such as "PUMP 43.53 BBLS PRE FLUSH" and "15% HCl MAIN ACID 21.3 BBLS". This demonstrated the robustness of the agentic workflow in handling technical language and numerical operations, even in cases prone to LLM misinterpretations like fluid type differentiation or concentration calculations.

Key observations include improved consistency due to schema enforcement, significant mitigation of common LLM errors (e.g., incorrect unit interpretation, misclassification of fluid types, erroneous percentage calculations), and accurate results even in ambiguous cases. By eliminating numerical hallucinations and ensuring accurate data aggregation, the framework consistently outperforms raw LLM approaches, which often yield inconsistent or incorrect calculations despite prompt tuning. This guided, tool-augmented approach enhances LLMs' contextual understanding while delegating mathematical fidelity to deterministic tools, providing a robust method for deriving accurate engineering insights from complex data. The framework is generalizable to other domains such as cementing, fracturing, and economic evaluation,

bridging natural language understanding with deterministic computation for scalable, trustworthy LLM deployment in real-world energy-sector applications.

Introduction

The advent of Large Language Models (LLMs) has ushered in a new era of artificial intelligence, profoundly impacting natural language processing, information retrieval, and contextual understanding across diverse industries. Their ability to generate coherent text, summarize vast documents, and even perform complex reasoning tasks has captivated researchers and practitioners alike. Within the engineering domain, particularly in sectors characterized by complex data, intricate processes, and critical safety considerations such as oil and gas, manufacturing, and civil infrastructure, the potential of LLMs to revolutionize data analysis, report generation, and operational insights is immense. Automating the extraction of key parameters, synthesizing information from different sources, and assisting in decision-making processes could significantly enhance efficiency, reduce manual effort, and unlock new levels of data-driven intelligence.

However, a fundamental challenge persists that significantly impedes the full realization of LLMs' potential in these precision-dependent fields: their inherent unreliability in mathematical operations and numerical consistency. Unlike symbolic computation systems or deterministic algorithms, LLMs are statistical models trained on vast datasets, learning patterns and correlations rather than logical rules or mathematical truths. This architectural characteristic often leads to what is colloquially termed "hallucination"—the generation of factually incorrect or nonsensical outputs—which is particularly pronounced when dealing with numerical computations, unit conversions, or precise data extraction. For example, an LLM might accurately identify a numerical value within a document but misinterpret its unit, transpose digits, fail to perform a simple arithmetic operation correctly, or even misinterpret contextual clues such as percentages (e.g., treating "15% HCl" as if it refers to 100% concentration) or fluid types (e.g., confusing "pre-flush" with "main acid volume, or misidentifying specialized fluids like Viscoelastic Diverting Acid).

In engineering, where precision is not merely a preference but a foundational requirement for safety, efficiency, and economic viability, such numerical inaccuracies are catastrophic. A miscalculated fluid volume in a well treatment, an incorrect pressure reading in a pipeline, or an erroneous depth interval in a geological report can lead to costly operational failures, environmental incidents, or erroneous design decisions. The current state of LLMs often necessitates extensive human oversight, manual verification, and recalculation of outputs, thereby negating much of the promised efficiency gains and introducing new avenues for human error. This "hallucination gap" thus creates a critical barrier to deploying LLMs for high-stakes analytical tasks, trapping their revolutionary capabilities within the realm of qualitative text processing, while leaving the numerical heavy lifting to traditional, often manual, methods. The problem is exacerbated by the sheer volume of unstructured data—reports, logs, schematics, and sensor data—that engineers must sift through daily, much of which contains vital numerical information embedded within complex technical prose. Furthermore, the non-deterministic nature of LLMs means that even with prompt tuning, results can vary from one run to the next, making it impossible to rely on their raw outputs for consistent, accurate calculations.

The framework directly addresses the pressing industry need for accurate, data-driven insights, enhancing operational efficiency and decision-making accuracy by providing a robust and trustworthy method for extracting and processing critical numerical data. An illustrative implementation, detailed in a subsequent section, substantiates the efficacy of this approach.

Proposed Solution: The Agentic AI Framework for Precision Engineering Operations

This section introduces an innovative agentic AI architecture that fundamentally transforms how Large Language Models (LLMs) are utilized for numerical operations in complex engineering contexts. Recognizing the inherent strengths of LLMs in natural language understanding and their limitations in deterministic computation, this framework orchestrates a multi-stage process that systematically guides LLMs, validates their outputs, and delegates precise calculations to external, reliable tools. This decomposition of tasks into a sequence of specialized, autonomous agents ensures consistent, dependable, and auditable high-precision results, effectively mitigating the "hallucination" problem while leveraging the contextual prowess of LLMs. The core principle behind this framework is the strategic separation of concerns: LLMs are employed for tasks where their capabilities are superior (e.g., semantic understanding, entity recognition in complex text), while deterministic, rule-based systems or traditional computational libraries are employed for tasks requiring absolute numerical precision and consistency (e.g., arithmetic, unit conversion, logical validation). This hybrid approach maximizes the utility of each component, creating a robust and trustworthy system for data extraction and processing in high-stakes engineering environments.

Core Architecture & Workflow

The framework operates through a modular, multi-agent pipeline, where each agent is responsible for a specific stage of the data processing workflow. This modularity not only enhances robustness and maintainability but also allows for independent optimization and scalability of individual components, as visually depicted in Fig 1.




Figure 1—Agentic AI Framework Architecture for Precision Engineering Operations

Document Agent

Function. This agent serves as the initial gateway for all incoming data, ingesting and preprocessing unstructured inputs such as PDF documents or plain text files. Its role is to convert varied inputs into a standardized, machine-readable format, often involving text extraction and basic cleaning operations.

Benefit. It ensures data consistency and prepares a clean input stream, laying the groundwork for accurate information extraction and reducing noise.

Extraction Agent

Function. Leveraging the advanced natural language understanding capabilities of an LLM, this agent identifies and extracts structured information from the pre-processed document. This involves carefully crafted prompts to guide the LLM to recognize specific entities, relationships, and numerical values

pertinent to the engineering domain. The LLM is guided to generate output directly conforming to a predefined structured format, such as JSON, simplifying subsequent validation and ensuring appropriate categorization of extracted data elements (e.g., differentiating fluid types or concentrations from volumes).

Benefit. This agent capitalizes on the LLM's superior contextual understanding and semantic parsing abilities, producing a higher initial quality of structured data ready for the next stage.

Validation Agent

Function. This agent acts as a critical "check-and-balance" layer. It utilizes predefined data schemas that explicitly define the expected structure, data types, and valid ranges for the extracted information. Any deviations or inconsistencies in the LLM's output are identified and flagged. Additional programmatic checks can be applied to ensure values are physically plausible and operationally sound, addressing potential LLM errors like misinterpretations of percentages or misclassifications.

Benefit. This agent directly mitigates common LLM errors such as incorrect unit interpretation, misaligned values, or numerical hallucinations. By enforcing strict data integrity, it prevents erroneous information from propagating downstream, significantly enhancing data trustworthiness.

Computation Agent

Function. This agent is responsible for executing deterministic calculations on the validated data. It entirely bypasses the LLM's numerical limitations by delegating mathematical fidelity to reliable, predictable computational tools. Operations performed include summing volumes, calculating effective concentrations, determining depth intervals, and counting stages. These calculations are executed by conventional programming constructs, ensuring absolute accuracy.

Benefit. This agent ensures absolute numerical precision and consistency, eliminating the risk of LLM-induced calculation errors and guaranteeing the accuracy and auditability of all numerical results, which is paramount for engineering applications.

Reporting Agent

Function. This final agent formats and outputs the processed results in a structured, usable, and easily consumable manner. This involves organizing computed values into structured tables or JSON, suitable for display, export, or seamless integration into business intelligence dashboards and other analytical platforms.

Benefit. This agent transforms raw, extracted, and computed figures into actionable insights, providing clean, structured data ready for immediate use in decision-making and operational workflows.

Technical Advantages & Innovation

This Agentic AI framework, as demonstrated by the provided illustrative implementation, provides several distinct technical advantages that position it as a significant innovation in the application of artificial intelligence to engineering challenges, particularly in the context of high-precision data handling.

Numerical Consistency through Deterministic Offloading

The most critical advantage of this framework is its ability to guarantee numerical consistency. This is achieved by performing all mathematical operations using reliable programming constructs and numerical libraries, rather than relying on the LLM for arithmetic. This strategic offloading, coupled with rigorous schema validation, ensures that numerical outputs are consistently dependable and auditable, directly addressing the core weakness of LLMs in mathematical precision and their non-deterministic calculation failures.

Active Mitigation of LLM "Hallucinations" and Errors

The multi-stage agentic workflow, particularly the Validation Agent and the structured output guidance provided to the LLM, actively corrects for common LLM errors. By enforcing a strict schema for the LLM's output and applying subsequent programmatic checks, it prevents common issues like misinterpretations of units, value transpositions, malformed responses, incorrect percentage interpretations, or misclassification of fluid types. This proactive enforcement significantly reduces the propagation of LLM "hallucinations" and misalignments, creating a more robust system.

Robustness in Handling Ambiguity and Complex Language

Through sophisticated prompt engineering and intelligent validation logic, the system can accurately process ambiguous cases and complex technical labeling common in unstructured engineering documents. The LLM is guided to interpret contextually nuanced phrases and extract the correct numerical and categorical information. This capability is inherent to powerful LLMs and is further refined by the framework's validation layers, ensuring logical coherence of the extracted data even when the original text is subtle.

Scalable and Generalizable Modular Architecture

The modular design of the framework, with distinct agents and the use of flexible, widely adopted tools, allows for easy adaptation and scalability. Each agent can be independently updated, optimized, or replaced with alternative technologies. This architecture demonstrates broad applicability across various unstructured engineering documents and calculation types simply by modifying the data schemas and the deterministic calculation logic, showcasing its versatility and potential for enterprise-wide deployment.

Enhanced Trustworthiness for High-Stakes AI Deployment

By integrating human-auditable, deterministic computation and explicit schema-based validation layers, the framework provides a robust and transparent approach. This significantly enhances the trustworthiness and utility of LLMs for real-world, high-stakes energy-sector applications. The explicit workflow makes the process auditable, contrasting with "black-box" LLM solutions and building confidence in AI-driven insights among engineers and stakeholders.

Unique Hybrid of Natural Language Processing and Deterministic Computation

This framework effectively illustrates a core innovation by having a powerful LLM perform intelligent natural language processing and entity extraction from unstructured documents, while conventional, reliable programming tools perform the precise, deterministic mathematical operations. This synergistic integration bridges a long-standing gap in AI applications, enabling automated analysis of complex engineering data with unprecedented accuracy and reliability, offering a solution superior to either component alone.

Business Value & Strategic Impact

The adoption and implementation of this Agentic AI framework offer significant strategic and operational benefits across the energy services sector. This innovation directly contributes to efficiency gains, risk mitigation, and the creation of new revenue opportunities for organizations.

Enhanced Operational Efficiency and Productivity

The framework automates the extraction, validation, and calculation of critical data from vast volumes of unstructured engineering reports. This drastically reduces the manual processing time currently required to parse documents, manually extract parameters, and perform calculations. Engineers and data analysts can reallocate their valuable time from tedious data handling to higher-value tasks such as analysis, interpretation, and strategic decision-making. This directly translates to increased productivity across various operational units, from well planning and execution to asset management and maintenance. For

example, rapidly processing hundreds of historical acid job reports can provide insights into optimal treatment parameters far more quickly than manual review.

Significant Risk Reduction and Improved Decision Quality

Minimizing the risk of costly errors arising from inaccurate or incomplete data is paramount in the energy industry. Manual data entry and human verification are prone to oversights, especially with large datasets and complex calculations. This Agentic AI framework, with its rigorous validation and deterministic computation layers, significantly reduces the probability of errors in critical engineering processes. Accurate data enables better-informed decisions, whether in optimizing acid jobs, designing well interventions, assessing reservoir performance, or predicting equipment failures. This leads to improved operational outcomes, enhanced safety, and reduced financial liabilities associated with faulty data.

New Service Offerings and Revenue Streams

The generalizable and scalable nature of this framework opens up possibilities for new digital solutions and services for service companies and operators. For example, offering automated analysis of client-provided historical operational data, generating compliance reports, or providing real-time data validation for ongoing projects could create new revenue streams. The framework's adaptability allows it to be customized for various client needs across different sub-segments of the energy industry, offering a competitive differentiator.

Strengthened Competitive Advantage through Innovation

Developing and deploying such an advanced, reliable AI framework demonstrates a commitment to cutting-edge innovation. This technological leadership creates a significant competitive advantage, differentiating organizations in the market by offering superior data processing capabilities. It showcases internal R&D prowess and establishes a high barrier to entry for competitors attempting to develop similar high-precision AI solutions for engineering data, fostering further innovation and reinforcing market leadership.

Accelerated Digital Transformation and AI Adoption

This framework serves as a practical, robust, and trustworthy pathway for deeper digital transformation and AI adoption within organizations and the broader industry. By demonstrating how LLMs can be reliably integrated into critical engineering workflows, it builds confidence in AI technologies. This successful implementation can serve as a blueprint for extending AI capabilities to other complex data challenges, fostering a culture of innovation and data-driven excellence throughout the organization. It bridges the gap between the theoretical promise of AI and its reliable real-world application.

Illustrative Use Case

The Agentic AI framework has been successfully tested and validated through a specific, illustrative implementation using real-world acid job reports. These reports typically contain a wealth of unstructured text detailing various operational parameters, fluid volumes, and treatment stages, often presenting significant challenges for automated parsing due to varied phrasing, implicit context, and inconsistent formatting of numerical data.

This implementation demonstrates the end-to-end workflow, beginning with the Document Agent ingesting a PDF file. The core of the extraction is handled by the Extraction Agent, which utilizes a powerful LLM. Strategic prompt design guides the LLM to extract specific parameters. This includes instructing the LLM to correctly separate numerical values (like "35.7 BBLS") from contextual descriptors (like "15% HCl" or "PRE FLUSH"), ensuring the values are associated with their correct fluid types and concentrations, and accurately identifying specialized fluid types like "VDA" (Viscoelastic Diverting Acid) and their associated volumes. The LLM is configured to produce structured output, which streamlines subsequent validation.

The Validation Agent then formally defines the precise structure and data types for the extracted information using a data schema. This programmatic validation ensures that extracted values conform to expected formats (e.g., numerical fields are indeed numbers, fluid types belong to predefined categories, and percentages are within valid ranges). This pre-validates the LLM's output and significantly mitigates common hallucination errors, ensuring data integrity before any computation, especially regarding misclassified fluid volumes or misinterpretations of percentages or specialized additives.

Upon successful extraction and validation, the Computation Agent takes over. Critical numerical calculations are performed using reliable, deterministic programming tools, entirely bypassing the LLM's limitations in arithmetic. For instance, the system accurately:

- Aggregates total volumes by summing values categorized into specific fluid types (e.g., pre-flush, main acid).
- Calculates effective acid volume by multiplying the main acid volume by its validated concentration.
- Calculates the total VDA (Viscoelastic Diverting Acid) driver by summing relevant fields.
- Determines the overall depth interval by finding the maximum and minimum depths from all stages and calculating their difference.
- Counts the total number of distinct stages identified.

These operations are deterministic and auditable, ensuring the numerical fidelity required for high-stakes engineering applications. Finally, the Reporting Agent formats these validated and computed results into a structured, usable output format. This output, containing summarized job values like total number of stages, aggregated acid volume, total VDA, and overall depth interval, provides clean, structured data ready for integration into dashboards, further analysis, or export to other enterprise systems.

Comparison: Raw LLM Output vs. Agentic Framework Output

To highlight the critical impact of the Agentic AI framework, a direct comparison was made between the output generated by a raw LLM (asked to extract and summarize numerical data without the agentic workflow) and the output from our proposed framework using the same acid job reports.

When a raw LLM was tasked with directly extracting numerical data and performing simple aggregations from the unstructured engineering reports, common issues observed included:

- **Numerical Hallucinations:** The LLM would occasionally generate numerical values that were not present in the original document or incorrectly combine digits.
- **Unit Misinterpretations/Omissions:** Units were often dropped, misinterpreted, or inconsistently attached to numerical values.
- **Inaccurate Aggregations & Percentage Misinterpretations:** The LLM frequently performed incorrect arithmetic when summing volumes, and often misinterpreted concentrations, leading to erroneous calculations (e.g., treating a percentage like "15%" as a direct multiplier rather than a fraction).
- **Fluid Type Misclassification:** The LLM often struggled to consistently differentiate between fluid types like "pre-flush," "main acid," or specialized additives like "VDA" (Viscoelastic Diverting Acid), leading to incorrect aggregation of volumes into their respective categories.
- **Contextual Misalignments:** Volumes or parameters might be correctly extracted but associated with the wrong fluid stage or depth range.
- **Unstructured/Inconsistent Output:** The format of the extracted data was highly variable, often requiring significant manual post-processing. Furthermore, due to the non-deterministic nature,

these errors could appear inconsistently across multiple runs, making debugging and reliance impossible.

In stark contrast, the Agentic AI framework consistently delivered highly accurate and reliable results. The sequential application of specialized agents effectively mitigated these issues:

- The **Extraction Agent**, guided by carefully designed prompts, accurately extracted numerical and contextual information, including identifying specialized fluid types like VDA, reducing initial errors.
- The **Validation Agent**, using strict data schemas, enforced type consistency and validated against predefined categories and ranges, catching and preventing misclassified fluid types or incorrect numerical interpretations.
- The **Computation Agent** ensured absolute mathematical fidelity by delegating all calculations to reliable, deterministic programming logic, completely eliminating arithmetic and percentage-related errors from the LLM.
- The **Reporting Agent** guaranteed consistent and structured output, ready for immediate use.

This comparison clearly demonstrates that while LLMs offer powerful natural language understanding, their integration into a robust agentic framework, complete with validation and deterministic computation, is essential for achieving the precision and trustworthiness required for critical engineering operations, especially when dealing with nuanced numerical interpretations and calculations. A detailed side-by-side comparison illustrating these differences is presented in [Table 1](#).

Table 1—Comparison of Raw LLM Output vs. Agentic AI Framework Output

Metric / Scenario	Raw LLM Output (Common Errors Observed)	Agentic AI Framework Output	Comment/Reason for Discrepancy Mitigation
Input Phrase Example	"15% HCl MAIN ACID 21.3 BBLS"	"15% HCl MAIN ACID 21.3 BBLS"	Consistent input for comparison.
Concentration Calculation	Incorrectly scales volume (e.g., $15 * 21.3$) or ignores percentage.	Extracts concentration (15%) and volume (21.3 BBLS) separately, then deterministically calculates effective acid.	LLM Computational/Interpretation Flaw: Raw LLM struggles with direct calculations and contextual percentage use. Framework separates interpretation and delegates calculation.
Fluid Type Classification	Incorrectly groups volumes (e.g., pre-flush counted as main acid, VDA misidentified)	Accurately categorizes and separates different fluid volumes (e.g., "pre_flush", "main acid", "VDA").	LLM Contextual Misinterpretation: Raw LLM fails to consistently distinguish fluid types. Framework uses schema validation and prompt guidance for precise categorization.
Total Aggregated Volume	Inconsistent or erroneous sums.	Accurate, consistent aggregated total volumes.	LLM Arithmetic/Non-Determinism: Raw LLM performs unreliable arithmetic, often varying. Framework's deterministic computation ensures precise totals.
Output Format	Highly variable, often unstructured.	Consistently structured (e.g., JSON, DataFrame).	Lack of Schema Enforcement: Raw LLM outputs are unpredictable. Framework enforces predefined schemas for consistent, usable data.
Overall Reliability	Low, necessitates extensive manual verification.	High, auditable, consistent results.	Inherent LLM Limitations: Raw LLMs are prone to numerical hallucinations and inconsistencies. The Agentic Framework's layered approach provides robustness and reliability.

Discussion and Limitations

While the Agentic AI framework offers significant advancements, its practical implementation comes with certain challenges and inherent limitations that warrant discussion:

Computational Cost and Latency

Leveraging powerful LLMs, especially for high-volume document processing, can incur substantial API costs and introduce latency in the overall workflow. Optimizing prompt efficiency and exploring local, smaller LLMs for certain tasks are ongoing considerations to manage these.

Prompt Engineering Complexity and Maintenance

Crafting highly precise and robust prompts for the Extraction Agent to handle the vast variability and nuance of engineering language requires significant domain expertise and iterative refinement. Maintaining these prompts and schemas as new document types emerge or as industry terminology evolves presents an ongoing management challenge, potentially requiring version control and automated testing.

Dependency on Input Quality

The framework's performance, particularly the initial text extraction by the Document Agent, is still dependent on the quality of the unstructured input. Poorly scanned documents, heavily annotated PDFs, or extremely ambiguous phrasing can still present difficulties that even sophisticated LLMs struggle with, potentially requiring manual intervention for very challenging cases.

Domain Specificity of Schemas

While the framework's architecture is generalizable, adapting it to new domains (e.g., cementing vs. fracturing) necessitates the development of new, highly tailored data schemas and deterministic calculation logic. This involves a non-trivial upfront effort in domain analysis and schema design for each new application area.

Handling True Ambiguity

Although the system significantly reduces misinterpretations by delegating deterministic tasks and validating outputs, genuinely ambiguous statements or incomplete information in source documents may persist. The framework is designed to flag these rather than hallucinate, but ultimate resolution might still require human expert review, particularly in critical high-stakes scenarios where automated decision-making isn't permissible.

Future Work

Building upon the successful implementation and the demonstrated advantages of this Agentic AI framework, several avenues for future work are envisioned to further enhance its capabilities and expand its applicability:

Expanded Domain and Document Types

While tested with acid job reports, future efforts will focus on implementing and validating the framework across a broader range of unstructured engineering documents, including cementing designs, fracturing reports, and complex reservoir evaluation summaries. This will involve developing new domain-specific schemas and expanding the library of deterministic computation tools.

Enhanced LLM Integration and Adaptivity

Investigating more advanced LLM techniques such as few-shot learning, fine-tuning on proprietary datasets, or incorporating self-correction mechanisms within the Extraction Agent could further improve parsing accuracy for highly specialized or ambiguous engineering jargon. Exploring real-time adaptive prompting based on validation feedback is also a promising direction.

Enhanced Knowledge Base Integration

Integrating the framework with existing knowledge databases, enterprise data warehouses, and semantic graphs to enrich the extracted data with contextual metadata, historical performance trends, and expert knowledge. This will enable more sophisticated analytical queries and deeper insights.

Integration with Enterprise Systems

Developing robust connectors for seamless integration with existing enterprise resource planning (ERP) systems, simulation software, and data visualization platforms (e.g., Power BI, Tableau). This would enable automated data flow from unstructured reports directly into operational workflows and decision-making dashboards.

Uncertainty Quantification and Confidence Scoring

For critical engineering applications, it's vital not just to provide an answer but also a measure of confidence. Future work will explore methods to quantify uncertainty in extracted values and provide confidence scores for the overall processing, especially for values at the edge of predefined ranges or from highly ambiguous text.

Multi-Modal Data Processing

Extending the Document Agent's capabilities to process multi-modal inputs, such as extracting information from embedded diagrams, graphs, or handwritten notes within engineering reports. This would require integrating advanced image processing and computer vision techniques.

Human-in-the-Loop Feedback Mechanisms

Implementing intuitive user interfaces that allow engineers to easily review flagged data, correct errors, and provide feedback. This feedback can then be used to iteratively improve the LLM prompts, validation schemas, and computational logic, fostering continuous learning and model refinement.

These next steps aim to evolve the Agentic AI framework into an even more comprehensive, intelligent, and seamlessly integrated solution for precision data processing across the energy industry.

Conclusion

The "Agentic AI: A Tool-Guided Approach to Achieving Precision in Engineering Numerical Operations" framework represents a significant advancement in the practical application of artificial intelligence within precision-dependent engineering domains. By ingeniously combining the strengths of Large Language Models for contextual understanding with the unwavering reliability of deterministic computation, this innovative architecture directly addresses the critical challenge of numerical inaccuracy and "hallucination" that has historically limited the deployment of LLMs in high-stakes environments, particularly the inconsistency and errors encountered with percentage interpretations and fluid type classifications, including specialized additives like Viscoelastic Diverting Acid.

Through its modular, multi-agent pipeline—comprising the Document, Extraction, Validation, Computation, and Reporting Agents—and validated by its successful application, the framework establishes a robust, auditable, and highly dependable system for extracting and processing critical numerical data from unstructured engineering documents. This application specifically demonstrates the power of using a powerful LLM constrained by data schemas for precise extraction, followed by deterministic calculations using reliable programming tools for numerical accuracy, thereby overcoming the inherent non-deterministic nature of raw LLM outputs.

This unique hybrid approach ensures numerical consistency, actively mitigates LLM errors, demonstrates robustness in handling complex linguistic ambiguities, and offers unparalleled scalability and generalizability across diverse engineering applications. The strategic adoption and potential

implementation of this framework offer substantial business value across the industry, including enhanced operational efficiency through automation, significant risk reduction by minimizing data errors, and the opening of new avenues for digital service offerings. Beyond these tangible benefits, it fosters a deeper level of trust in AI deployment within critical industry sectors, paving the way for more widespread and impactful integration of intelligent systems. As demonstrated through its successful application in analyzing acid job reports and its superior performance compared to raw LLM outputs, this Agentic AI framework is not merely a theoretical construct but a practical, deployable solution that bridges the gap between cutting-edge AI capabilities and the stringent demands of engineering precision, marking a pivotal step toward truly trustworthy and transformative AI in the energy industry.

References

- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. 2017. Attention Is All You Need. *Advances in Neural Information Processing Systems*, **30**. <https://doi.org/10.48550/arXiv.1706.03762>
- Yao, S., Cui, J., Li, Z., Luo, J., Xiao, J., Wu, X., et al 2023. ReAct: Synergizing Reasoning and Acting in Language Models. *Proceedings of the International Conference on Learning Representations (ICLR)*. <https://doi.org/10.48550/arXiv.2210.03629>
- Schick, T., Schick, M., and Schütze, H. 2023. Toolformer: Language Models Can Teach Themselves to Use Tools. *Proceedings of the International Conference on Learning Representations (ICLR)*. <https://doi.org/10.48550/arXiv.2302.04761>
- Park, J. S., O'Shaughnessy, J. C., Huang, K. H., Chan, M., Navi, A., Hwang, J., et al 2023. Generative Agents: Interactive Simulacra of Human Behavior. *Proceedings of the ACM Symposium on User Interface Software and Technology (UIST)*. <https://doi.org/10.1145/3586183.3606763>