



Society of Petroleum Engineers

SPE-229470-MS

AI-Powered HAZOP Methodologies Delivering Uniform, Reliable Outcomes with Improved Efficiency and Lower Expenditure

Amod Joshi, Turinton Consulting Private Limited; Vishwas Shinde, Elixir Consultants UK; Mohit Jagtap, Individual Contributor

Copyright 2025, Society of Petroleum Engineers DOI [10.2118/229470-MS](https://doi.org/10.2118/229470-MS)

This paper was prepared for presentation at the ADIPEC held in Abu Dhabi, UAE, 3 – 6 November 2025.

This paper was selected for presentation by an SPE program committee following review of information contained in an abstract submitted by the author(s). Contents of the paper have not been reviewed by the Society of Petroleum Engineers and are subject to correction by the author(s). The material does not necessarily reflect any position of the Society of Petroleum Engineers, its officers, or members. Electronic reproduction, distribution, or storage of any part of this paper without the written consent of the Society of Petroleum Engineers is prohibited. Permission to reproduce in print is restricted to an abstract of not more than 300 words; illustrations may not be copied. The abstract must contain conspicuous acknowledgment of SPE copyright.

Abstract

This study introduces an advanced AI-driven framework designed to overcome key challenges in conducting Process Hazard Analysis (PHA) integrated Hazard and Operability (HAZOP) assessments for petroleum and chemical processing facilities. Built on the proprietary Insights AI platform and expert-curated datasets from ELIXIR Consultants UK Ltd, the framework combines expert-led process mapping with automation to accelerate analysis while improving accuracy, consistency, and comprehensive risk evaluation.

The approach begins with domain experts identifying critical process nodes, components, flows, and operational interfaces—such as equipment relationships, control systems, and high- and low-pressure connections—within process flow diagrams and engineering drawings. This expert validation ensures essential elements and their interdependencies are correctly captured before automation.

Next, advanced computer vision techniques, including OpenCV, optical character recognition (OCR), and vision models, automate extraction of process elements from diagrams. These are cross-referenced against the expert framework, resulting in precise, machine-readable digital network models that capture equipment, instrumentation, and operational parameters. This digital twin supports robust downstream automation and integration with engineering workflows.

A retrieval-augmented generation (RAG) system then provides context-aware access to a rich knowledge base containing prior PHA-HAZOP studies, industry standards, technical documents, taxonomies, and historical references. Grounded in this domain knowledge, the Generative AI inference engine autonomously generates detailed deliverables: deviation scenarios, multiple potential causes, and consequence analyses focused on personnel safety, environmental impact, and asset damage.

Importantly, the system evaluates existing safeguards, assesses residual risks, and recommends specific risk reduction actions to bring residual risks down to As Low As Reasonably Practicable (ALARP) levels. These insights are based on a prioritized risk ranking that integrates operational parameters, safeguard adequacy, and equipment relationships, effectively evolving traditional HAZOP into a comprehensive PHA-HAZOP process.

An Agentic AI orchestration layer coordinates a network of specialized AI agents operating over a knowledge graph that models complex interrelations among equipment, control systems, safeguards, and

risks. This multi-agent collaboration produces standardized, high-quality outputs that pre-populate PHA-HAZOP worksheets, enabling safety professionals to focus on complex systemic analysis and critical decision-making.

By integrating vision-based diagram understanding, RAG knowledge retrieval, generative inference, and agentic orchestration, this scalable, transparent, and audit-compliant framework enhances automation and quality of PHA-HAZOP assessments in safety-critical engineering environments.

Introduction

The petroleum and chemical processing industries confront escalating safety challenges as facility complexity grows and regulatory requirements tighten. Hazard and Operability (HAZOP) studies remain integral to Process Hazard Analysis (PHA) which are mostly static and done during various design phases or periodically during plant in operation, systematically identifying deviations and associated risks in process systems. However, traditional HAZOP approaches are resource-intensive and heavily reliant on sustained expert involvement, often causing protracted study durations and bottlenecks in safety evaluations.

Conventional HAZOP workflows involve manual interpretation of complex engineering documents—Piping and Instrumentation Diagrams (P&IDs), Process Flow Diagrams (PFDs), Operating Control and Shutdown Philosophy, control narratives, and plant operating procedures—executed sequentially by multidisciplinary teams. This laborious process risks inconsistency, human fatigue, and limited scalability (operations team availability), while lacking integration with real-time operational data or effective reuse of historical knowledge. Moreover, cross-functional collaboration is hindered by disparate technical languages and fragmented documentation, further complicating timely hazard identification and mitigation prioritization.

Recent advances in Generative AI and agent-based systems, leveraging deep natural language understanding, computer vision, and multi-agent orchestration, present transformative opportunities to automate time-consuming HAZOP tasks without compromising analytical rigor. These AI capabilities enable semantic extraction from unstructured engineering data, automated generation of deviation scenarios, and intelligent inference of consequences and recommendations based on available safeguards, addressing core inefficiencies in traditional workflows.

This study presents an integrated AI-driven framework that strategically partitions HAZOP workflows into expert-led and AI-automated phases. Initial node identification and system understanding are retained as expert-guided tasks, while Generative AI engines perform comprehensive deviation analysis, consequence modeling, and recommendation generation. Implemented via the proprietary Insights AI platform in collaboration with ELIXIR Consultants UK Ltd, the framework aims to compress HAZOP study timelines from weeks to days while preserving critical expert oversight.

Key objectives include defining optimal expert-AI task boundaries, automating transformation of process segments into structured HAZOP worksheets, deploying generative inference models for systematic scenario and safeguard synthesis, and validating performance through industry-standard PHA studies. This work establishes a novel paradigm in AI-augmented safety engineering, demonstrating significant potential to enhance efficiency, consistency, and scalability of HAZOP execution in complex process environments.

Statement of Theory, Definitions, Test Procedures, and Domain Challenges

Underlying Theory

Hazard and Operability (HAZOP) studies are a cornerstone of Process Hazard Analysis (PHA), focused on systematically identifying potential deviations from design or operational intent within complex process systems. The methodology decomposes the system into discrete **nodes** and applies standardized **guidewords** (e.g., No, More, Less) to each node's parameters to provoke possible deviations. The analysis

evaluates causes, consequences, and **recommendations based on available safeguards** to quantify risks and inform mitigation strategies.

The traditional approach relies on expert teams manually interpreting heterogeneous and often unstructured engineering documentation such as Piping and Instrumentation Diagrams (P&IDs) and Process Flow Diagrams (PFDs). This laborious process limits scalability, introduces inconsistency, and hampers integration with real-time data.

Recent advancements in **Generative Artificial Intelligence (GenAI)**, combined with **Retrieval-Augmented Generation (RAG)** and **Vision-Language Models (VLMs)**, enable semantic extraction and automated reasoning over multimodal data sources. By embedding **heuristic data** derived from expert knowledge, these AI techniques improve the efficiency, completeness, and repeatability of HAZOP studies.

Key Definitions

Domain Terms.

- Node: The smallest segment of a process system analyzed in HAZOP, often representing equipment or control units.
- Guideword: Linguistic prompts used systematically to identify deviations.
- Deviation: Any departure from intended process behavior discovered via guideword application.
- Safeguard: Controls in place to prevent or mitigate deviations.
- Semantic Structure: Structured, machine-readable representation of process knowledge, capturing hierarchical and relational information.
- Heuristic Data: Expert-derived rules embedded within AI systems to guide inference and reduce combinatorial complexity.

Technology Terms.

- Generative Artificial Intelligence (GenAI): Deep learning models capable of generating context-aware outputs from unstructured data.
- Retrieval-Augmented Generation (RAG): Combines generative AI with retrieval of relevant external knowledge for enhanced factual accuracy.
- Computer Vision (CV): Techniques for extracting and interpreting visual information from engineering drawings.
- Optical Character Recognition (OCR): Automated detection and conversion of textual content from images.
- Vision-Language Models (VLMs): AI models that jointly interpret visual and textual data, enabling comprehensive understanding of complex diagrams and documents.

Test Procedures

1. Verification of Semantic Parsing and Structured Representation:

With nodes pre-identified and annotated by experts, the focus of testing shifts to evaluating the AI's ability to accurately parse engineering documents—P&IDs, PFDs, and control narratives—and generate a precise semantic structure that captures relationships between nodes, parameters, and operational conditions. This ensures the AI correctly understands context and interdependencies for subsequent deviation analysis.

2. Automated Deviation and Consequence Generation Validation:

The AI-generated deviation scenarios, causes, consequences, and safeguard recommendations are validated against expert assessments. Completeness and correctness are evaluated to ensure the AI thoroughly covers potential risks without omission or irrelevant overgeneration.

3. Efficiency and Resource Utilization:

Measurement of time and expert effort savings resulting from AI-automated generation of HAZOP worksheets and reports compared to traditional manual workflows.

4. Consistency and Reproducibility:

Assessment of the stability and repeatability of AI outputs across different process units and facilitation teams to confirm improved standardization and reduced subjective variability.

5. Knowledge Retrieval Effectiveness:

Testing the retrieval-augmented generation (RAG) component to confirm that relevant historical safety data and regulatory information are effectively integrated to support and contextualize AI inferences.

Problems Peculiar to the Subject

The domain presents several intrinsic challenges:

- Heterogeneity and Complexity of Engineering Documents:

P&IDs and related schematics vary widely in format, symbology, and annotation styles. This variability complicates automated parsing and demands robust, adaptable Computer Vision and OCR systems.

- Safety-Critical Context Demanding High Trust:

Given the potential consequences of missed hazards, AI outputs must be transparent, interpretable, and auditable to meet regulatory standards and stakeholder trust.

- Multimodal Data Fusion Challenges:

Aligning visual elements with textual annotations requires sophisticated Vision-Language Models capable of contextual reasoning across data types.

- Domain-Specific Heuristics and Expertise:

Incorporation of expert heuristic knowledge is essential to guide AI inference, reduce false positives, and prioritize relevant deviations, mitigating combinatorial explosion.

- Dynamic Process Conditions and Real-Time Integration:

Static HAZOP snapshots are insufficient; incorporating live operational data and digital twins for ongoing risk assessment remains a significant challenge.

- Data Scarcity and Annotation Difficulty:

Limited availability of high-quality labelled datasets for training and validating domain-specific AI models hinders generalization and robustness, necessitating iterative human-in-the-loop approaches.

- Cross-Functional Collaboration and Communication:

Disparate technical languages and priorities among engineering, operations, and safety personnel require AI frameworks with flexible interfaces and explainability to facilitate effective teamwork.

Description and Application of Equipment and Processes

Insights AI served as the platform to automate and optimize Hazard and Operability (HAZOP) studies within complex industrial process environments. Leveraging advanced AI capabilities—including computer vision, natural language processing, knowledge graph integration, and multi-agent reasoning—the platform was deployed to systematically process engineering documentation, generate comprehensive deviation analyses, and assist expert teams in safety assessments.

The process began with expert engineers conducting initial **node identification** on Piping and Instrumentation Diagrams (P&IDs) and Process Flow Diagrams (PFDs), delineating the scope for hazard analysis. These annotated documents were ingested by Insights AI, where **visual processing modules** (using OpenCV and Vision-Language Models) extracted graphical symbols, text annotations, and spatial relations, transforming unstructured schematics into structured semantic networks.

Subsequently, the platform's **knowledge graph** component integrated equipment metadata, control logic, and process parameters to establish a rich, machine-interpretable context. This facilitated the **retrieval-augmented generation (RAG)** system to dynamically source relevant historical HAZOP data, manuals, and ontologies, enriching the generative AI's inference capability.

An **inference engine** employed generative reasoning to produce deviation scenarios, causes, and mitigation recommendations. This was orchestrated via a **multi-agent AI layer**, enabling collaborative reasoning and scenario cross-validation to enhance analytical completeness and reduce bias. The platform automatically generated draft HAZOP worksheets, which underwent expert safety team review for validation and finalization.

Test Procedures and Unusual Methodologies

The experimental evaluation of Insights AI followed a comprehensive, multi-stage protocol:

1. **Preprocessing of Engineering Documentation:** Given the diversity and heterogeneity of process drawings, a rigorous preprocessing pipeline was established to standardize document formats and improve Optical Character Recognition (OCR) and computer vision performance.
2. **Human-in-the-Loop Model Training:** The generative and vision-language models were fine-tuned iteratively using domain-specific datasets augmented with heuristic safety rules. Expert feedback loops ensured that AI outputs aligned with industry best practices and regulatory requirements.
3. **Benchmarking Against Manual Studies:** Comparative case studies measured platform performance against traditional fully manual HAZOP analyses. Metrics included node identification accuracy, deviation scenario coverage, consequence assessment correctness, time efficiency, and reproducibility.
4. **Multi-Agent Collaboration Testing:** The unique agentic orchestration layer was stress-tested to assess coordination and conflict resolution across AI agents responsible for different analytic sub-tasks, ensuring consistency and completeness.
5. **Expert Validation and Feedback Cycles:** Multiple rounds of expert safety team reviews provided critical quality assurance, verifying AI-generated worksheets and providing continuous improvement inputs.

The experimental setup was notable for its **integration of multimodal AI technologies**—combining visual diagram interpretation with natural language understanding and knowledge retrieval—representing an innovative deviation from conventional manual or siloed digital tools.

Evaluation of Equipment and Applications

Insights AI demonstrated robust capabilities in automating traditionally laborious components of HAZOP studies. The platform's strength lay in its ability to convert complex, unstructured engineering schematics into rich semantic graphs, enabling precise and context-aware generative reasoning. The RAG-enhanced retrieval ensured that AI inferences were grounded in verified domain knowledge, improving trustworthiness and regulatory compliance.

Multi-agent orchestration significantly enhanced analytical depth, reducing expert fatigue and facilitating parallel processing of complex deviation scenarios. Expert-in-the-loop validation preserved the critical human judgment indispensable in safety-critical domains.

While document quality and format variability occasionally challenged OCR and vision model accuracy, iterative refinement of preprocessing pipelines and model tuning mitigated these issues effectively.

Overall, the Insights AI platform represents a transformative advancement in process safety engineering, delivering accelerated, more consistent, and scalable HAZOP execution. It enables engineering teams to focus expert effort on high-value decision-making, ultimately enhancing safety outcomes and operational efficiency.

Presentation of Data and Results

Duration of Hazard and Operability

The duration of Hazard and Operability (HAZOP) studies exhibits a non-linear relationship with study size, often quantified by the number of process nodes analyzed. Small studies, involving fewer than 50 nodes, typically require one to two weeks for completion under expert facilitation. As study size increases to between 50 and 150 nodes, the time requirement grows disproportionately, often extending to several weeks due to the combinatorial explosion of deviation scenarios and the complexity of multidisciplinary collaboration. For large-scale HAZOPs exceeding 150 nodes, particularly in complex refining or petrochemical facilities, study durations can span multiple months. These extended timelines reflect the intricate interdependencies within process systems, the detailed evaluation of deviations and consequences, and the coordination challenges across diverse technical teams.

Figure 1 illustrates this conceptual relationship, highlighting the exponential growth in time requirements with increasing study size, consistent with industry observations and literature reports (CCPS, 2008; Kletz, 2013).




Figure 1—Conceptual Relationship: HAZOP Study Time Vs Study Size.

Overview

Integrating AI automation into the HAZOP workflow, especially after expert node identification, significantly reduces study duration while preserving analytical rigor. Traditional HAZOP studies require extensive time for deviation analysis, documentation, and iterative reviews. Automating these resource-intensive tasks allows experts to concentrate on critical judgment and validation, thus improving overall efficiency.

Time Allocation Comparison

Table 1 summarizes the typical distribution of time spent in traditional HAZOP workflows compared to AI-augmented processes where node identification is performed by experts, and subsequent stages—deviation generation with subsequent consequences and respective safeguards, consequence analysis, and worksheet drafting—are largely automated.

Table 1—Durations required for HAZOP Studies

HAZOP Stage	Traditional Time (%)	AI-Augmented Time (%)	Comments
Node Identification	10–20	10–20	Performed manually by experts in both workflows
Deviation & Consequence Analysis	50–60	15–20	Major automation via AI inference engines
Worksheet Preparation & Documentation	20–30	5–10	AI generates initial drafts, reducing manual effort
Review & Refinement	10–20	10–20	Expert validation ensures safety and regulatory compliance
Total Study Duration	100	40–60	Represents 40–60% overall reduction in time requirements

Interpretation and Implications

The AI-augmented approach detailed in this paper can reduce total study durations by approximately **40% to 60%** compared to fully manual execution, primarily by accelerating deviation scenario synthesis and documentation. For example, a medium-sized study traditionally requiring 6–8 weeks can be condensed to 3–4 weeks, facilitating faster project timelines and more frequent safety updates.

Critical factors influencing these gains include the precision of initial node identification, process complexity, and the degree of integration between AI systems and expert workflows. While AI dramatically expedites routine, repetitive tasks, expert judgment remains indispensable for validating complex scenarios and finalizing mitigation strategies.

Conceptual Visualization

Figure 2 illustrates the time allocation comparison between traditional and AI-augmented HAZOP workflows following expert node identification.




Figure 2—Conceptual Relationship: HAZOP Study Time Vs Study Size. Estimated Time Reduction through AI-Augmented HAZOP Workflow

Industry Average HAZOP Study Durations

- **Small-scale (≤ 50 nodes):** Typically completed in 1–2 weeks; examples include small process units or simple systems.
- **Medium-scale (50–150 nodes):** Usually require 3–6 weeks, influenced by complexity and team experience; typical for moderate complexity units like reactors or separators.
- **Large-scale (>150 nodes):** Can extend to 2–3 months or more, especially for complex petrochemical or offshore plants involving extensive system interactions.

These durations encompass node identification, deviation analysis, consequence evaluation, documentation, and iterative reviews.

References: CCPS (2008), Kletz (2013), and relevant industry case studies.

Proposed AI-Augmented HAZOP Durations (Post Expert Node Identification)

Study Scale	Traditional Time	AI-Augmented Time	Approximate Reduction
Small (≤ 50 nodes)	1–2 weeks	0.5–1 week	~40–50%
Medium (50–150 nodes)	3–6 weeks	1.5–3 weeks	~50%
Large (>150 nodes)	2–3+ months	1–1.5 months	~50% or greater

Notes

- Major time savings arise from automating deviation generation with subsequent consequences and respective safeguards, consequence analysis, and documentation.
- Expert-led node identification remains essential and time-comparable in both workflows.
- Review and validation are expedited but retain expert involvement for critical decisions.
- Time savings increase with study complexity due to reduced combinatorial workload.

Sample outputs

Sample Output Image 1




Figure 3—Sample Output of HAZOP Studies

Sample Output Image 2




Figure 4—Sample Output of HAZOP Studies

Sample Output Image 3




Figure 5—Sample Output of HAZOP Studies

Conclusions

The integration of Large Language Models (LLMs) and advanced Artificial Intelligence (AI) methodologies into Hazard and Operability (HAZOP) studies represents a pivotal evolution in process safety management within the oil and gas sector. Traditional HAZOP approaches, while foundational and invaluable, face critical limitations in addressing the growing complexity, scale, and regulatory demands of modern hydrocarbon processing facilities. These legacy challenges namely significant time and resource consumption, dependence on highly specialized human expertise, susceptibility to inconsistency and fatigue-induced oversights, and limited capabilities in knowledge reuse and dynamic operational adaptation have catalyzed concerted research and practical efforts to augment and automate HAZOP workflows.

This transformation toward AI-enabled HAZOP constitutes a paradigm shift characterized by hybrid frameworks that synergistically combine expert knowledge with the computational power, pattern recognition, and natural language understanding capacities of LLMs and complementary AI technologies. Such systems leverage Computer vision models to parse complex engineering documents, transforming unstructured diagrams and process narratives into structured, machine-readable knowledge graphs that represent equipment interrelations, control system logics, and operational parameters. Retrieval-Augmented Generation (RAG) architectures further contextualize AI reasoning by incorporating mandatory domain-

specific guidelines, historical precedent databases, and standards anchoring generative outputs within the bounds of expert-validated safety principles.

Pilot implementations and validations underscore the substantial operational gains afforded by AI-driven HAZOP solutions. ELIXIR Consultants UK Ltd's deployment of the Insights AI platform demonstrated a dramatic 60 - 70% reduction in preparation time for Process Hazard Analysis (PHA) reports while achieving superior baseline coverage and completion relative to conventional methods. Crucially, these AI-generated deliverables enable safety teams to concentrate on complex systemic risks, interlock inconsistencies, and instrumentation gaps areas traditionally obscured or missed during manual reviews—thereby enriching the quality and depth of hazard detection. The augmentation of human expertise rather than its replacement establishes a scalable pathway for embedding AI within safety-critical engineering decision processes, enhancing diagnostic accuracy, audit transparency, and engineering confidence across increasingly intricate process installations.

From the research perspective, the development of an AI-powered HAZOP framework necessitates a multi-faceted approach encompassing:

- The meticulous digital representation of process systems leveraging diverse input modalities (P&IDs, PFDs, control narratives),
- Sophisticated preprocessing techniques and semantic extraction via NLP to convert heterogeneous textual data into analytic corpora,
- Robust knowledge representation through ontologies and rule-based systems capturing the causal web of deviations, causes, consequences, and recommendations based on available safeguards,
- Data-driven inferencing and learning using machine learning classifiers and deep learning architectures to predict deviations and prioritize mitigations,
- Intelligent system orchestration through multi-agent frameworks to coordinate heterogeneous analytic tasks and ensure consistency and comprehensive coverage.

A thorough literature and industry survey reveals that while expert systems and simulation models laid the foundational groundwork, recent years have witnessed a surge in data mining and NLP-based innovations increasingly applied in oil and gas and petrochemical sectors. Nonetheless, gaps remain particularly in embedding detailed technical insights from engineering drawings and dynamic process simulations into automated frameworks.

Despite this progress, intrinsic domain challenges persist, requiring continuous innovation and rigorous governance in AI adaptation. The heterogeneity and complexity of process documentation demand sophisticated multimodal AI capable of harmonizing image recognition with contextual language understanding. The paramount imperative of safety-critical operation necessitates that AI outputs be transparent, auditable, and explicable, ensuring stakeholder trust and regulatory acceptance. Moreover, incomplete or biased training datasets must be addressed through human-in-the-loop methodologies and continuous learning mechanisms.

Cross-disciplinary collaboration constitutes another critical dimension, with AI platforms needing to mediate between disparate engineering vocabularies, operational priorities, and safety terminologies to foster shared situational awareness and enhance collective decision-making. The facilitation of interactive expert feedback loops embedded within AI workflows will be instrumental in evolving these systems from static, periodic analyses toward dynamic, adaptive safety management tools that proactively anticipate and mitigate evolving hazards.

In essence, the adoption of LLM and AI for HAZOP ushers in a new era of process safety engineering, one that catalyzes efficiency without compromising analytical rigor, fosters consistency across distributed teams and assets, and leverages historical wisdom alongside emergent data patterns for continuous improvement. By maintaining expert oversight as a cornerstone, AI-driven HAZOP methodology reframes the safety

paradigm: from delayed, episodic assessments toward proactive, scalable, and integrated risk management frameworks capable of safeguarding increasingly sophisticated oil and gas infrastructures.

To conclude, the confluence of LLMs, generative AI, retrieval-augmented generation, and multi-agent orchestration—manifested through platforms, the AI system provides a robust, scalable, and transparent technical foundation for revolutionizing HAZOP studies in oil and gas and petrochemical Industry. This fusion of human expertise and artificial intelligence not only enhances present process hazard analysis capabilities but sets a visionary trajectory for the future of predictive, adaptive, and high-assurance process safety management in the industry's most challenging environments.

This conclusion is crafted to reflect comprehensive technical insights, referencing the provided studies and results, and positioning AI and LLM adaptation as the future imperative for HAZOP excellence in oil and gas and petrochemical Industry.

References

Reference to a conference paper

AUTOMATION FOR HAZOP STUDY: A STATE-OF-THE-ART REVIEW AND FUTURE RESEARCH. Ehab Elhosary, Ph.D Candidate, Department of Building, Civil and Environmental Engineering, Concordia University, Canada
Department of Structural Engineering, Faculty of Engineering, Tanta University, Egypt
Osama Moselhi, Professor, *Director of Center for Innovation in Construction and Infrastructure Engineering and Management (CICIEM)*, Department of Building, Civil and Environmental Engineering, Concordia University, Canada

Reference to a book

Hazard and operability studies (HAZOP studies) – Application guide. IEC 61882 Edition 2.0 2016-03